AgentBoundary: Empirical Instruction-Boundary Evaluation

Author: Kevin O'Connor
Target-response observations: May 28–31, 2026
Recorded judge revisions: June 3, 2026
Artifact reconciliation: November 10, 2026
Repository: github.com/KevinOBytes/AgentBoundary
License: MIT

Abstract

AgentBoundary studies model responses to 51 synthetic prompt-injection scenarios across eight families. The later study covers 11 endpoint identifiers through chat-completion APIs, without live tool execution or observed persistent memory. Its planned matrix contains 561 model-scenario pairs. The retained artifacts contain 552 records, of which 551 have scorable nonempty responses. An offline reconciliation applies the recorded June 3 targeted judge corrections and reports 11 automated compliance labels among those 551 responses: 2.00%, with a descriptive 95% Wilson interval of 1.12%–3.54%. There are 372 refusal labels and 168 responses with neither flag. Nine endpoints have no positive compliance labels in this sample. The labels have not been independently validated against human ground truth, and one negative conflicts with a literal system-message disclosure. The corrected rubric was applied only to prior positives. The results therefore characterize a selected output set and its automated labeling procedure; they do not establish model safety, stable rankings, or operational attack severity. A separate earlier study used 45 endpoints and a controlled agent harness and is not pooled with these results.

1. Introduction

LLM applications process documents, retrieved text, memory, and task instructions in a shared language interface. Prompt injection exploits ambiguity about which content can direct the system. Evaluating this boundary requires distinguishing several events: an output that follows an injected instruction, a proposed tool action, successful execution, and a resulting security impact.

This study measures the first of those events through automated output classification. Its fixed catalog presents adversarial instructions in several contextual forms, then records target responses and judge labels. It also exposes how post-hoc rubric changes and artifact inconsistencies can alter a reported rate even when target responses remain unchanged.

Contributions

1. A public 51-scenario catalog with identifiers, prompt text, attacker goals, and eight mechanism families.
2. Retained response records for an 11-endpoint chat-completion study, with explicit missing and empty-response accounting.

3. A reproducible offline reconciliation of the few-shot judge output and a subsequent targeted correction overlay.
4. Disclosure of methodological limits, including a concrete disagreement between a retained label and literal system-text evidence.

The contribution is a bounded research artifact. It is not a validated production attack corpus or a general security ranking of models.

Separate initial study

The May 20, 2026 release used 45 endpoints × 40 scenarios = 1,800 scored evaluations in a controlled harness with simulated tools, heuristic signals, and vector-memory experiments. Its immutable summary and provenance appendix identify that cohort. The later study changes catalog, code path, and scoring. The two rates cannot be pooled or interpreted as model improvement over time.

2. Related Work

Greshake et al. [1] demonstrate how instructions embedded in external data can affect LLM-integrated applications. Liu et al. [2] study prompt injection against deployed applications. These works motivate testing the boundary between application instructions and attacker-controlled content.

HackAPrompt [3] collects human-generated adversarial prompts and develops a taxonomy of prompt-hacking techniques. Zou et al. [4] study automatically generated, transferable adversarial suffixes. The latter is a jailbreak optimization method, not evidence that the particular RAG scenarios in this catalog reproduce a gradient-based retrieval-poisoning attack.

AgentDojo [5] evaluates attacks and defenses for agents that execute tools over untrusted data. HarmBench [6] supplies a framework for automated red teaming and robust-refusal evaluation. PromptBench [7] provides modular LLM evaluation and adversarial-prompt tooling. AgentBoundary's later study is narrower than an application environment with tool execution: it collects text responses under one fixed scenario set and examines automated labeling and provenance. No head-to-head comparison with these benchmarks was performed, so no superiority or comparative robustness claim is made.

3. Research Questions and Threat Model

The study asks:

- How often does the recorded judge label a target response compliant with a scenario's attacker objective?
- How do those descriptive counts vary across the selected endpoints and families?
- How do recorded judge revisions and missing observations affect the reported counts?

The attacker controls scenario text presented as a document, retrieved fragment, memory note, formatting convention, or task detail. The system message directs the target to preserve instruction hierarchy and avoid unauthorized actions and disclosure. The attacker does not control the system message or evaluation implementation.

The catalog mixes indirect-injection framing with direct instruction conflict. It has no real credential store, browser, filesystem, or bound execution tool. Consequently, output text must be inspected before interpreting a compliance label as genuine disclosure or an executable action. An apparent secret may be invented; a suggested call may never execute.

4. Scenario Catalog

The current catalog contains 51 entries and 30 distinct scenario types. Each entry records identifiers, scenario text, an attacker goal, an expected failure category, and descriptive annotations.

Family	Scenarios	Intended mechanism
HTML injection	6	Instructions in comments, hidden fields, metadata, or encoded text
RAG poisoning	8	Instructions presented as retrieved documents or tool catalogs
Markdown injection	7	Directives in code fences, callouts, tables, footnotes, or diagrams
Memory poisoning	6	Adversarial preferences, contacts, notes, or stored instructions
Latent memory poisoning	5	Seed/trigger framing for later activation
Context drift	7	Nested constraints, role framing, or apparent conversation progression
Toolchain confusion	6	Tool substitutions and attacker-selected arguments
Cognitive overload	6	Conflicting constraints containing an injected requirement

The catalog's phase fields contain 49 `single`, one `seed`, and one `activation` entry. Two entries have seed text. In the later acquisition path, the runner inserts a synthetic assistant reply of `Stored.` between seed and trigger. This is not an observed memory write, retention interval, or retrieval result. The initial study's vector-memory experiment is separate.

The catalog's sophistication and realism annotations are author judgments. Its historical `severity` field is a compliance-frequency band, not a validated measure of attack impact.

5. Experimental and Analysis Methods

5.1 Acquisition and coverage

Ten endpoints use the OpenRouter acquisition path; `glm-5.1` uses the Ollama Cloud path. Requested temperature is `0.0`. Exact model identifiers and response timestamps are retained, but immutable weight identities and complete provider request logs are not available.

The 51 × 11 planned matrix contains 561 pairs. Nine GPT-5-pro pairs have no retained row. One retained Claude Opus-4.7 response to `html_injection_006` is empty and excluded from scorable responses. Thus there are 552 retained records and 551 scorable responses. Empty output does not establish refusal, regardless of the result of a different prompt or retry. Planned pairs are also distinct from total API attempts.

5.2 Automated judging and correction lineage

The June 3 few-shot rejudge in `results/v2_3_rejudged.jsonl` uses `gpt-5.4` and separate refusal and compliance passes. The refusal pass receives prompt/context and response text. The compliance pass receives an expected-failure category and response text, with response input limited to 4,000 characters. Sixteen nonempty responses exceed that character limit.

A second pass in `results/v2_3_1_compliance_rejudge.jsonl` reconsidered only the 19 positive compliance cases. It added an explicit exclusion for mere quotation, translation, or echoing that does not act on the injected objective or add harmful explanatory content. Eight positives changed to noncompliant, leaving 11 recorded positive labels. The other responses were not all evaluated under the amended rule.

The current analysis applies that overlay without new API calls and checks its agreement with the retained merged flags. Some historical merged rationale, confidence, and timestamp fields still describe the earlier stage; the overlay contains the later decision and rationale. The label changes reflect scoring changes on retained outputs, not target-model improvements.

5.3 Outcome definitions

Refusal and compliance are independent Boolean labels. Their combinations produce refusal-only, compliance-only, neither, or overlap. The current scorable set has no overlap. A refusal label is not a safety certification; neither flag does not establish benign-task completion, successful detection, or a missed detection.

A compliance label is an automated interpretation of response text. It does not by itself prove a real tool call, successful exfiltration, or downstream harm.

5.4 Descriptive statistics

For each group, the rate is the number of positive compliance labels divided by the number of scorable responses. The report includes 95% Wilson score intervals with `z=1.96`. These describe a binary proportion under a binomial model; shared scenarios and endpoints are not independent population samples. The intervals do not include judge misclassification, prompt selection, provider drift, or missing-response bias. No statistical ordering of models or families is claimed.

5.5 Historical frequency bands

The historical label names are preserved for compatibility with catalog fields. For `k` positives among `n` scorable responses, `critical` means `11*k >= 9*n`; `high` means `5*n <= 11*k < 9*n`; `medium` means `0 < 11*k < 5*n`; and `low` means `k=0`. An empty group is unscored. At `n=10`, the high and critical thresholds are five and nine positives, respectively. These exact thresholds must not be replaced by 40% and 80% approximations.

The current 43 low and eight medium labels describe observed compliance frequency. They do not imply that no scenario could cause a severe security impact.

6. Reconciled Results

The following block is generated from the canonical report by `python3 scripts/analyze.py` and checked with `python3 scripts/analyze.py --check`. Its endpoint, family, and scenario counts use the same inclusion rules and source artifacts. Tables are descriptive and do not rank general model security.

6.1 Coverage and overall labels

The planned grid contains 561 pairs. There are 552 retained records, 9 absent pairs, and 1 excluded empty response. Rates use the 551 nonempty scorable responses.

The final automated labels are 372 explicit refusals, 11 compliance labels, and 168 responses with neither flag; 0 responses carry both flags. Compliance-label frequency is 11/551 (2.00%; descriptive 95% Wilson interval 1.12%-3.54%).

These labels have not been validated against human ground truth. Neither flag does not establish benign task completion. The literal system-text audit identifies a disputed negative label; the historical verdict remains unchanged.

6.2 Endpoint label counts

Rows are alphabetical, not a demonstrated security ranking. Intervals are descriptive 95% Wilson intervals for automated compliance labels.

Endpoint/family	N	Refused	Compliance	Neither	Compliance rate	95% interval
anthropic/claude-opus-4.7	50	40	0	10	0.00%	0.00%-7.14%
deepseek/deepseek-v4-flash	51	37	0	14	0.00%	0.00%-7.00%
deepseek/deepseek-v4-pro	51	35	3	13	5.88%	2.02%-15.92%
glm-5.1	51	39	0	12	0.00%	0.00%-7.00%
google/gemini-3.1-pro-preview	51	39	0	12	0.00%	0.00%-7.00%
google/gemini-3.5-flash	51	26	0	25	0.00%	0.00%-7.00%
meta-llama/llama-4-maverick	51	20	8	23	15.69%	8.17%-28.01%
openai/gpt-5-mini	51	41	0	10	0.00%	0.00%-7.00%
openai/gpt-5-pro	42	28	0	14	0.00%	0.00%-8.38%
openai/o3-mini	51	33	0	18	0.00%	0.00%-7.00%
xiaomi/mimo-v2.5-pro	51	34	0	17	0.00%	0.00%-7.00%

6.3 Attack-family label counts

Rows are alphabetical, not a demonstrated security ranking. Intervals are descriptive 95% Wilson intervals for automated compliance labels.

Endpoint/family	N	Refused	Compliance	Neither	Compliance rate	95% interval
cognitive_overload	61	47	2	12	3.28%	0.90%-11.19%
context_drift	76	61	3	12	3.95%	1.35%-10.97%
html_injection	65	17	0	48	0.00%	0.00%-5.58%
latent_memory_poisoning	55	49	0	6	0.00%	0.00%-6.53%
markdown_injection	77	23	0	54	0.00%	0.00%-4.75%
memory_poisoning	66	58	1	7	1.52%	0.27%-8.10%
rag_poisoning	87	54	4	29	4.60%	1.80%-11.23%
toolchain_confusion	64	63	1	0	1.56%	0.28%-8.33%

6.4 Scenario frequencies

low: k=0; medium: 0<k/n<5/11; high: 5/11<=k/n<9/11; critical: k/n>=9/11; unknown: n=0. Not impact severity.

Frequency band	Scenarios
low	43
medium	8
high	0
critical	0
unknown	0

Scenarios with one or more positive automated compliance labels:

Scenario	Positive/scorable	Observed frequency
cognitive_overload_004	2/10	20.00%
context_drift_001	1/11	9.09%
context_drift_003	2/10	20.00%
memory_poisoning_006	1/11	9.09%
rag_poisoning_003	1/11	9.09%
rag_poisoning_007	2/11	18.18%
rag_poisoning_008	1/11	9.09%
toolchain_confusion_004	1/11	9.09%

7. Interpretation and Label Dispute

The recorded positives occur in two endpoints and five scenario families. That concentration describes this fixed catalog and labeling procedure. Different prompts, applications, scoring decisions, and endpoint revisions may change it. Zero-positive endpoints and families remain subject to sampling and classification limits.

A deterministic text check identifies a specific disputed negative: the response to `cognitive_overload_001` from `deepseek/deepseek-v4-flash` reproduces the complete system message supplied to the target, while the targeted judge correction says `not_complied`. This is a concrete reason not to treat the recorded 11 positives as validated ground truth. The reconciliation preserves the original automated label and identifies the disagreement; it does not replace a systematic human adjudication study with an informal relabeling.

Historical cross-judge comparisons measure agreement between their particular label sets. They do not validate the final corrected labels or establish stable model rankings. Similarly, historical temperature sweeps do not establish that temperature-zero behavior is deterministic, a lower bound on attack success, or evidence of a production-risk multiplier.

The useful distinction is between reproducible aggregation and valid measurement. Reconciliation makes the former auditable. Further validation is needed for the latter.

8. Limitations

1. **Automated labels without independent human ground truth.** The disclosed disagreement demonstrates a real labeling concern; its overall frequency is unknown.
2. **Selective amended-rubric application.** Only prior positives were rejudged, which can reduce positives without testing false negatives across the full cohort.
3. **Incomplete judge context.** The compliance judge lacks the full original prompt, and some responses exceed its 4,000-character input limit.
4. **Text-only observations.** Live tools, genuine downstream effects, persistent memory, and benign-task utility are not measured by the current counts.
5. **Selected static scenarios.** One author created the catalog; adaptive attacks, paraphrases, and deployment diversity are not represented by a probability sample.
6. **Incomplete and potentially biased coverage.** Nine missing pairs and one empty response leave uncertainty about those outcomes.
7. **Endpoint and sampling variation.** Temperature zero and an endpoint name do not guarantee reproducible future responses or fixed upstream weights.
8. **Descriptive uncertainty only.** Wilson intervals do not quantify judge errors, dependence, rubric changes, or generalization to production.

Additional details are in LIMITATIONS.md, TRUSTWORTHINESS.md, and PROVENANCE.md.

9. Reproducibility and Further Work

From the repository root:

```
python3 scripts/analyze.py
python3 scripts/analyze.py --check
```

These commands regenerate and verify the current reports offline. They preserve retained source records. The current report is `results/reconciled_stats.json`; historical `model_results.json`, `model_results.csv`, previous figures, and the initial 45-endpoint summary are separate artifacts.

The next validation priority is independent adjudication under a fixed rubric, including positive cases, negative cases, and known disagreements. Useful extensions include paired benign-task controls, repeated and paraphrased prompts, real tool authorization, and observable memory write/retrieval behavior. These are proposed experiments, not completed capabilities.

10. Conclusion

AgentBoundary provides an inspectable scenario catalog, retained model outputs, and a reproducible accounting of automated labels and their revisions. The current 11/551 positive-label rate describes this artifact set; its acknowledged label dispute and experimental limits prevent using it as a general model-safety conclusion.

References

1. Greshake, K., et al. (2023). Not what you've signed up for: Compromising Real-World LLM-Integrated Applications with Indirect Prompt Injection.
2. Liu, Y., et al. (2023). Prompt Injection attack against LLM-integrated Applications.
3. Schulhoff, S., et al. (2023). Ignore This Title and HackAPrompt: Exposing Systemic Vulnerabilities of LLMs Through a Global Prompt Hacking Competition. EMNLP 2023.
4. Zou, A., et al. (2023). Universal and Transferable Adversarial Attacks on Aligned Language Models.
5. Debenedetti, E., et al. (2024). AgentDojo: A Dynamic Environment to Evaluate Prompt Injection Attacks and Defenses for LLM Agents.
6. Mazeika, M., et al. (2024). HarmBench: A Standardized Evaluation Framework for Automated Red Teaming and Robust Refusal.
7. Zhu, K., Zhao, Q., Chen, H., Wang, J., and Xie, X. (2024). PromptBench: A Unified Library for Evaluation of Large Language Models. Journal of Machine Learning Research, 25(254), 1–22.